Computer Architecture Lab 3 Report

Fabian Wüthrich

November 2020

The task of this lab is to extend Ramulator to evaluate two memory scheduling policies: ATLAS and BLISS. Before the actual implementation of the new policies, the scheduler was refactored using the strategy design pattern to make the code easier to extend. The refactoring removes many nested if-statements and moves fields, that were only relevant to certain policies, to the respective policy class.

Task 1: Implementing ATLAS

For implementing ATLAS, we need to track for each thread the attained service (AS) in the current quantum and the overall AS (TotalAS) up to the current quantum. Ramulator executes only one thread per core, so `coreid` is used to identify a thread.

Each memory controller (see `Controller.h`) tracks the AS of a thread during a quantum in the vector `attained_service`. When a request is issued, the corresponding vector entry is incremented. The counter is enclosed in a if-statement so the counter is only incremented on the first command of a request (see line 417).

ATLAS requires a meta-controller to coordinate multiple memory controllers. The meta-controller is implemented in the `Memory` class because the class holds references to all memory controllers. At the end of a quantum, the meta-controller fetches and accumulates the AS values for each thread (see line 295-302 in `Memory.h`). Then, the meta-controller recalculates the TotalAS for each thread and passes the new value to all memory-controllers (see line 304-313).

The actual policy is implemented in `Scheduler.h` by overriding the `compare()` function. The scheduler compares the arrival time of a request to the clock count of the memory controller to figure out if the request has been outstanding for more than T cycles (see line 229-233) . If both requests are below the threshold T, the scheduler queries the TotalAS value for each request from the memory controller and picks the request with the lowest value (see line 235-241). If both AS values are the same, the scheduler prioritizes row hits and then oldest requests first.

Task 2: Implementing BLISS

BLISS requires even less machinery than ATLAS and is implemented in `Controller.h`. The vector `blacklisted` keeps track of the blacklisted threads. Two variables (`last_request_coreid` and `request_served_counter`) are used to identify malicious threads.

Similar to ATLAS the `compare()` function is used to pick the best request. The scheduler returns the request whose thread is not blacklisted or defaults to row hits/oldest first if a decision based on the blacklist is not possible.

Task 3: Evaluation

The evaluation compares ATLAS and BLISS to three scheduling policies (FCFS, FRFCFS, FRFCFS_Cap) which are already implemented in Ramulator. For each scheduling policy, we simulated the following workloads:

- Workload 1: `HLLL` (four-core)

- Workload 2: `HHLL` (four-core)

- Workload 3: `HHHH` (four-core)

- Workload 4: `HHHHHHHH` (eight-core)

where `H` stands for a workload with high memory intensity and `L` is a workload with low memory intensity. Each simulation runs until every core retires 20 million instructions. We evaluated the instruction throughput (IT) and max. slowdown (MS), as defined in the lab sheet, for each workload using each policy.

Figure 1 shows the IT for each scheduling policy using different workloads. For the two mixed workloads (`HLLL`/`HHLL`), all policies have a similar IT. For the two intensive workloads (`HHHH`/`HHHHHHHH`), ATLAS has the lowest IT of all policies and BLISS outperforms the other policies for the eight core workload.

Figure 2 shows the MS for each policy. Similarly to the IT, we cannot observe substantial differences for the two mixed workloads. For the `HHHH` workload the MS increases for FCFS and ATLAS. For the eight core workload, only BLISS is able to keep the MS low i.e. guarantee fairness among the cores.

We expected that ATLAS outperforms the other policies for the mixed workloads because threads with less attained service i.e. the L workloads are prioritized. Our simulation uses a system with just one memory controller and four cores, whereas ATLAS was designed for multiple memory controllers in a many core system, which explains the poor performance of ATLAS.

The performance of ATLAS degrades significantly for the memory intensive workloads (`HHHH`/`HHHHHHHH`). This has primarily two reasons. First, the wrong workloads are used to show the benefits of ATLAS. The simulation runs the same trace accessing the same addresses on four (eight) cores simultaneously. Such workloads are rare in practice and schedulers which prioritize row buffer locality, benefit from such workloads. Second, ATLAS is not a fair scheduling algorithm, so instead of dividing the DRAM equally among all cores, requests of one core could be stalled for a long time degrading throughput. This unfairness is also reflected in the MS for ATLAS. For the eight core workload, some cores are over 4x slower in a multi-core setting as when running alone on the system.

Figure 1: Instruction Throughput

BLISS on the other hand is a fair scheduling algorithm. The MS for BLISS running the eight-core workload is fairly low, whereas the unfairness of the other scheduling algorithms increases. For the same reason, BLISS has a better IT for the eight-core workload than the other scheduling policies.

Figure 2: Max. Slowdown